Hands on with Apache MXNet

On the Amazon Deep Learning AMI

Ric Harvey - Technical Developer Evangelist
@ric__harvey

AI and ML

- Artificial Intelligence: design software applications which exhibit human-like behavior, e.g. speech, natural language processing, reasoning or intuition
- Machine Learning: teach machines to learn without being explicitly programmed
- Deep Learning: using neural networks, teach machines to learn from complex data where features cannot be explicitly expressed

Deep learning

Deep Learning is a subfield of machine learning concerned with algorithms inspired by the structure and function of the brain called artificial neural networks.

Data is passed through multiple non-linear transformations to generate a prediction, models use a cascade of multiple layers of nonlinear processing units for feature extraction and transformation. Each successive layer uses the output from the previous layer as input.

Objective: Learn the parameters of the transformations that minimize a cost function

The advent of Deep Learning

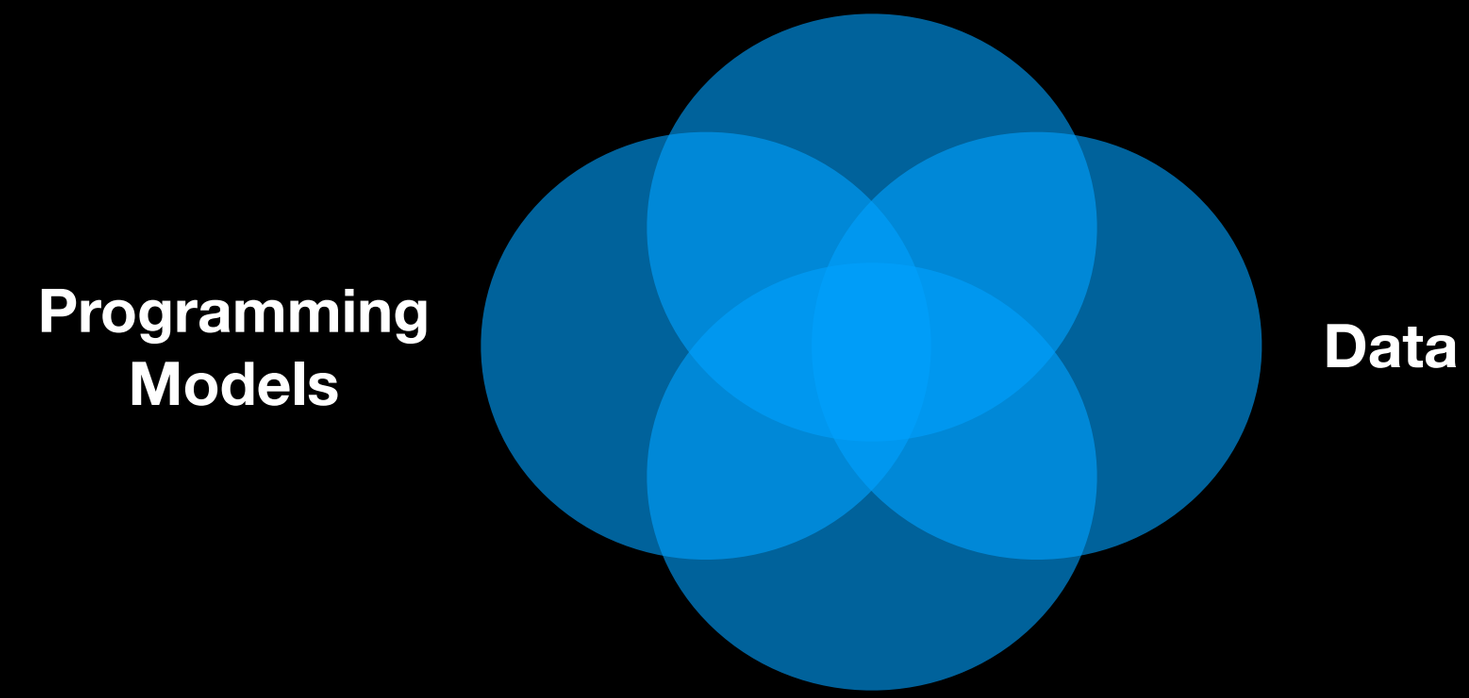

Uses of Deep Learning

Image understanding

- Expedia have 10M+ images from 300K+ hotels
- Images boost the conversation around a hotel
- So having the best images matter
- They used Keras and EC2 GPU instances and fine tuned a retrained model

Speech recognition

- Key word trigger

Natural language processing

- Intents….

Autonomy

- Object Segmentation
- Last June, tuSimple drove an autonomous truck for 200 miles from Yuma, AZ to San Diego

Customers Running AI on AWS

NETFLIX · The Washington Post · realnetworks · Expedia · Stanford

图森 tuSimple · Happy Snap · zmags · HubSpot · OhioHealth · RNIB

duolingo · Pinterest · iTranslate · intel · slack · NVIDIA

WOLFRAM · Carnegie Mellon · American Heart Association · MapBox · Zillow · NASA

SAMSUNG SmartThings · FINRA · GoAnimate · University of Washington · Capital One · C-SPAN

And many more…

aws

How does Deep Learning work?

Human Neuron

Inputs

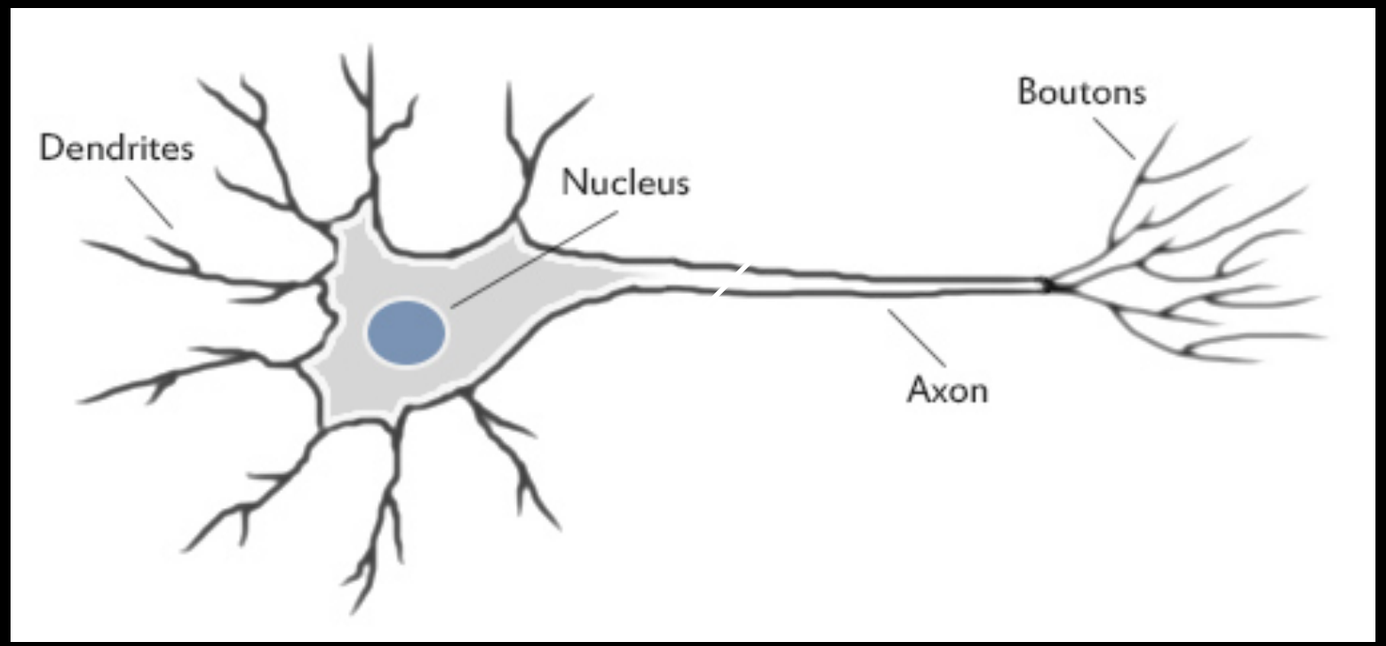

Output

Artificial Neuron and Perception

- Input
  - Vector training of data set (x)
- Output
  - Liner function of input
- Nonlinearity
  - Transformation of output into value range
- Training
  - Learn the weights and bias (b) by minimize loss

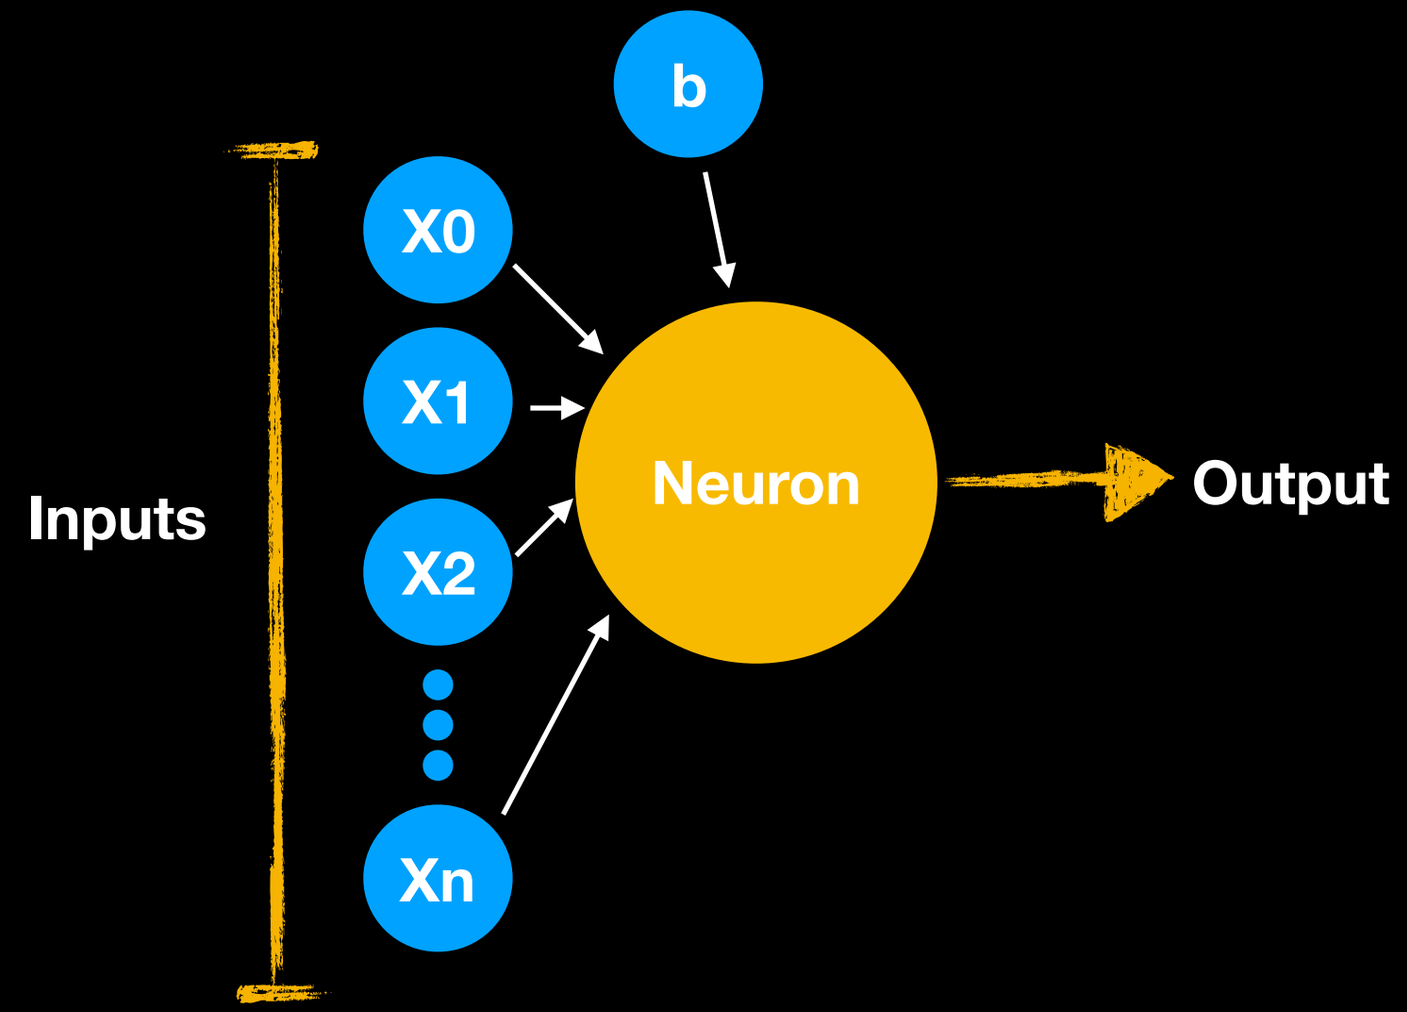

$$f(x) = \sigma\left(\langle w, x \rangle + b\right)$$

Models of Neural Networks

Lots of types, feed forward, recurrent neural network, radial based function, regulatory feedback and so on….

- **Convolution Neural Network(CNN)**
  - **Feedforward network**
  - **Inspired by the visual cortex and responds to stimuli in a restricted area**
  - **Good for image processing**
- **Long Short Term Memory(LSTM)**
  - **Propagates data forward and also backwards from later stages to earlier**
  - **LSTM out performed every other RNN model**

Apache MXNet

aws

https://mxnet.apache.org/

aws

Features

- **Flexible:** Supports both imperative and symbolic programming
- **Portable:** Runs on CPUs or GPUs, on clusters, servers, desktops, or mobile phones
- **Multiple Languages:** C++, Python, R, Scala, Julia, Matlab, Javascript, and Perl
- **Distributed on Cloud:** Supports distributed training on multiple CPU/GPU machines
- **Performance Optimized:** Optimized C++ backend engine parallelizes both I/O and computation
- **Broad Model Support:** CNN, RNN/LSTM

MXNet architecture

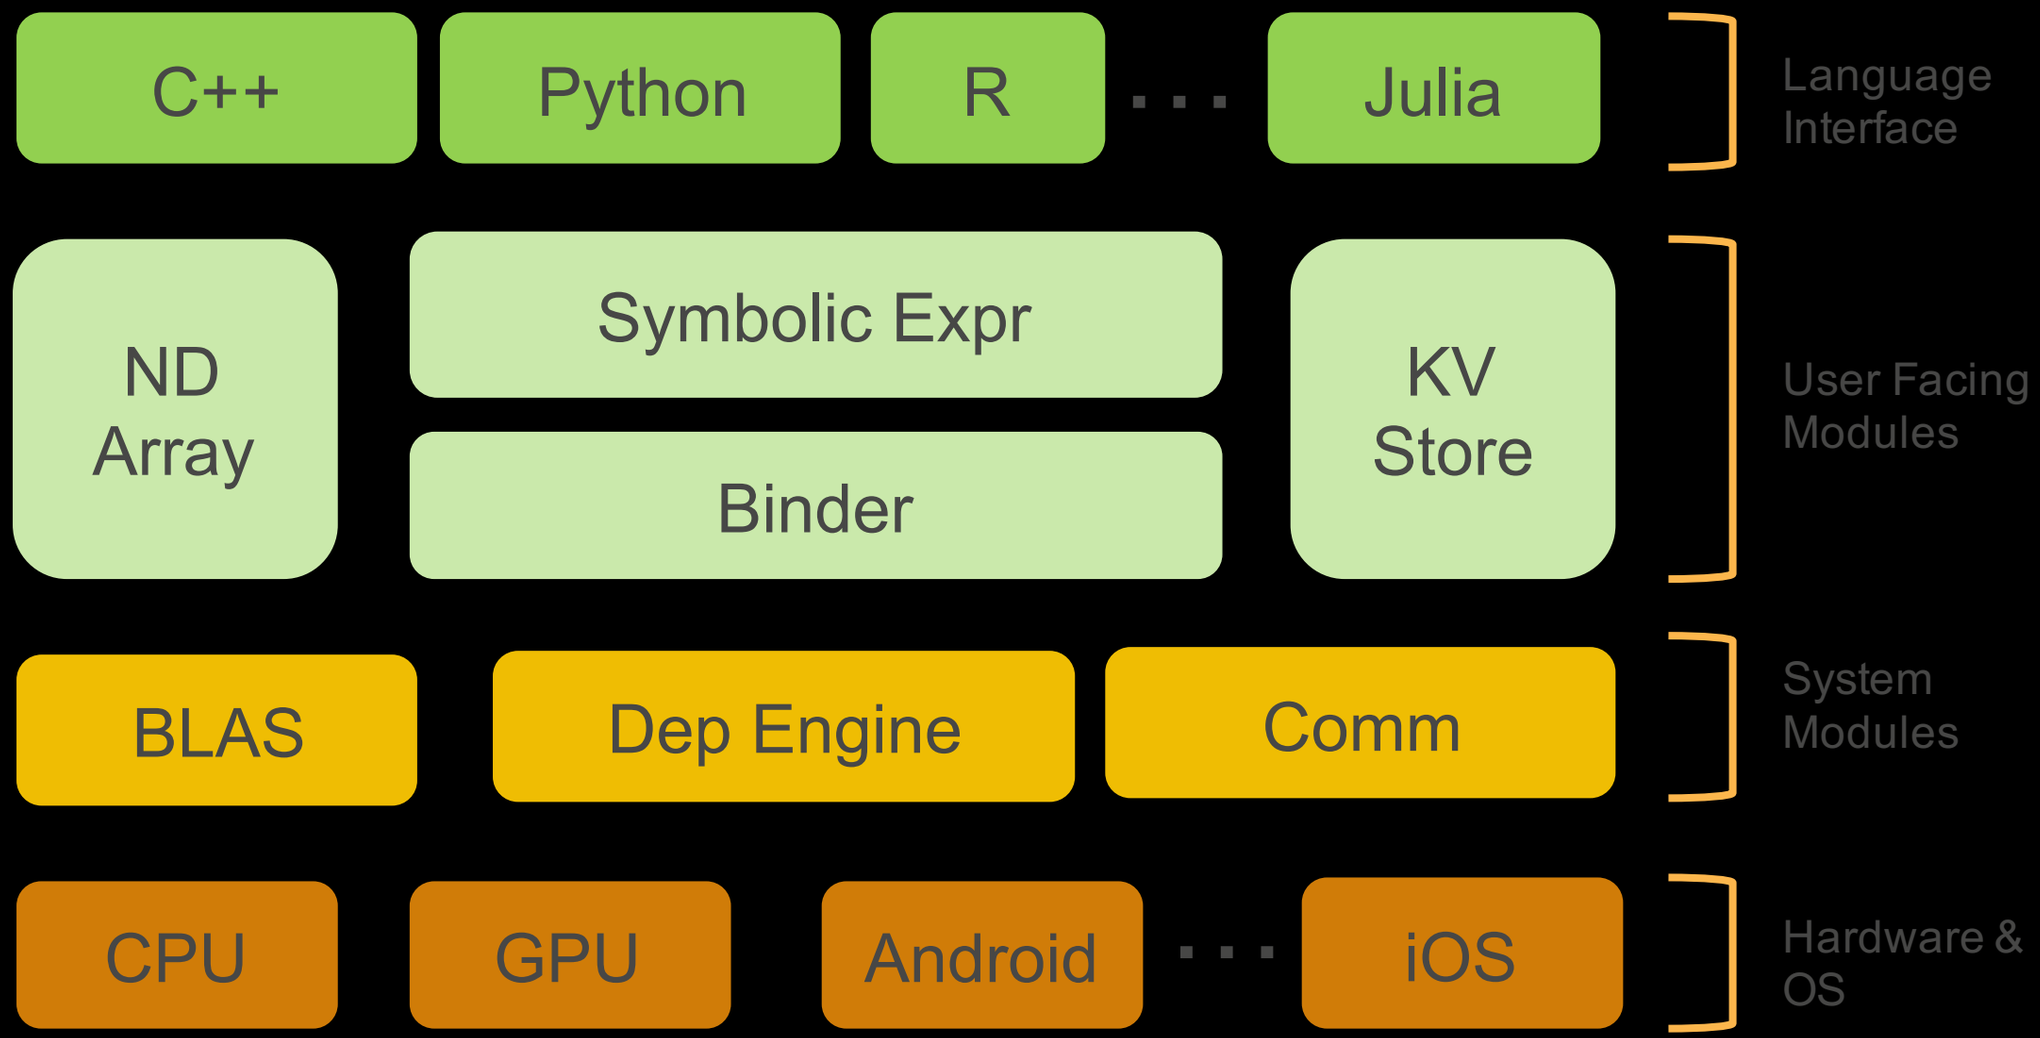

MXNet API Components

- **NDArray:** Provides imperative tensor operations
- **Symbol:** Provides neural network graph and auto-differentiation
- **RNN Cell:** Tools for building RNN symbolic graph
- **Module:** Provides interface for performing computation with Symbol
- **Data Loading:** Provides iterators for reading data
- **Metric:** Evaluation metric to evaluate performance of trained model

Model training flow in MXNet

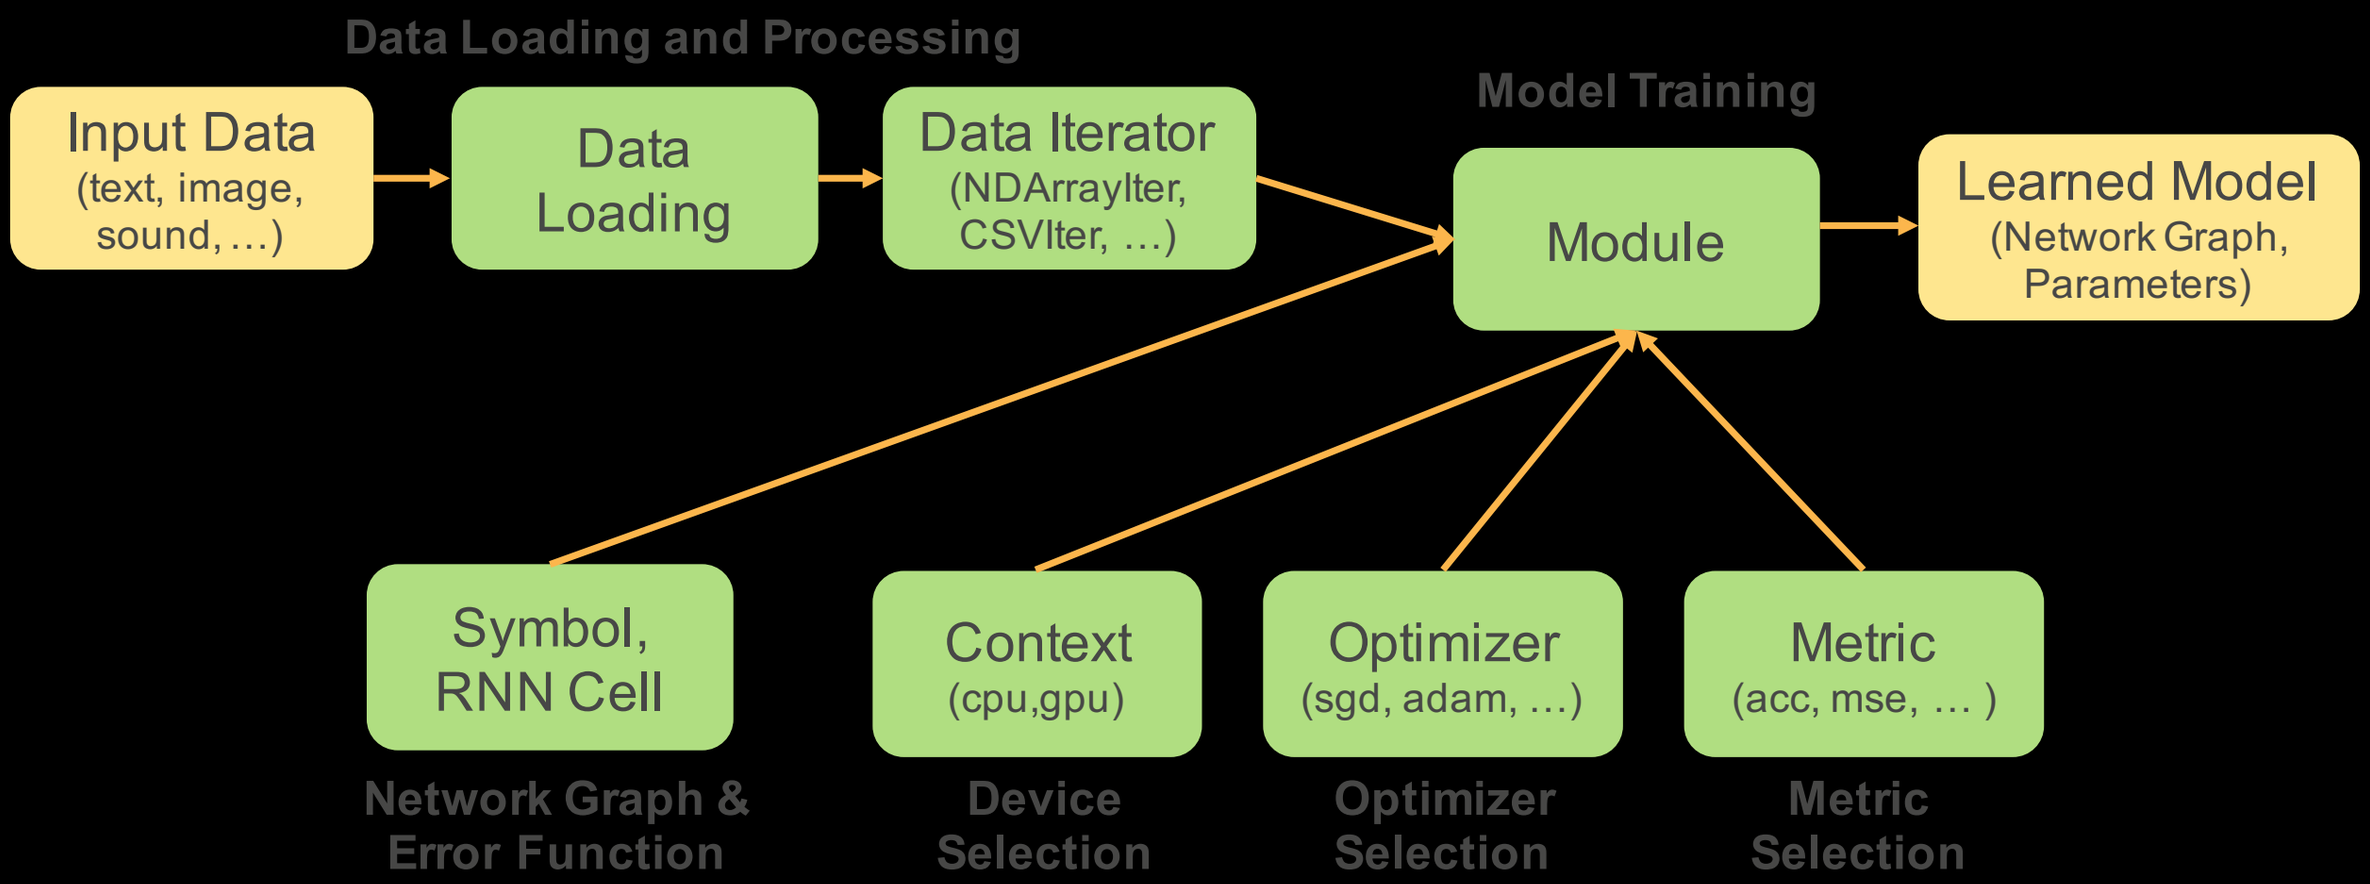

Workshop

MXNet with a pre-trained model

Exercise

Use the Amazon Deep Learning AMI to identify whats in these pictures and compare pre-trained models.

Conclusion

Model Comparison

How much memory does it use?

We can take an educated guess by looking at the size of the parameters file

- VGG16: 528MB (about 140 million parameters)
- ResNet-152: 230MB (about 60 million parameters)
- Inception v3: 43MB (about 25 million parameters)

How fast can it predict?

This is more difficult and can depend on batch size but in our example, lets look at the average over a few calls

*** VGG16
Predicted in 0.30 millisecond
*** ResNet-152
Predicted in 0.90 millisecond
*** Inception v3
Predicted in 0.40 millisecond

Summary

In these tests

- **ResNet-152 has the best accuracy of all three networks (by far) but it's also 2–3 times slower.**
- **VGG16 is the fastest—due its small number of layers?—but it has the highest memory usage and the worst accuracy.**
- **Inception v3 is almost as fast, while delivering better accuracy and the most conservative memory usage. This last point makes it a good candidate for constrained environments.**